



REPT: Reverse Debugging of Failures in Deployed Software


Harshitha Nagapudi - hnagapud@ttu.edu

Texas Tech University

OUTLINE

- + Existing Problems
- + REPT as a Solution
- + How it Works
- + Evaluation
- + Limitations
- + Future Work
- + Conclusion







REPT: Reverse Execution with Processor Trace





REPT: Reverse Execution with Processor Trace

- Online hardware tracing (e.g., Intel Processor Trace)
 - Log the control flow with timestamps
 - Low runtime overhead (1 – 5%)
 - *No data!*
- Offline binary analysis
 - Recovers data flow from the control flow




REPT Data Recovery

- Single-threaded execution reconstruction
- Multi-threaded execution reconstruction

Core Dump

```
10101011010101010101011001101010010011110110110
110101011011010101101010101010110011010100100
1111011011011010101101101010110001010101010110
0110101001001111011011010101101101010110101010
101010110011010100100111101101101101011011010
10110100101010101011001101010010011110110110110
101011011010101101010101010110011010100100111
10110110101010110110101011010101010101100110
101001001111011011011010110110101011010110101
010101010110011010100100111101101101101011011
01010110101010101010110011010100100111101101101
01011010101010101011001101010010011110110110110
101011011010101101010101010110011010100100111
101101101101010110110101010101010101010101010
1100110101001001111011011011010101011010110101
010101010110011010100100111101101101010110101
0101010101100110101001001111011011010101101101
01011010101010101011001101010010011110110110110
1010110110101011010101010101010101011001101010
010011110110110110101011011010101101010101010
1100110101001001111011011011010101101010110101
01101010101010101100110101001001111011011011010
```

Instruction Sequence



Execution History

= ?



How to recover overwritten states


```
lea rbx, [g]
```

```
mov rax, 1
```

```
add rax, [rbx]
```

```
mov [rbx], rax
```

```
xor rbx, rbx
```

lea rbx, [g]

mov rax, 1

add rax, [rbx]

mov [rbx], rax

➡ xor rbx, rbx

rax=?, rbx=?, [g]=3

rax=?, rbx=?, [g]=3

rax=?, rbx=?, [g]=3

rax=3, rbx=?, [g]=3

rax=3, rbx=?, [g]=3

rax=3, rbx=0, [g]=3

WRONG!

➔ `lea rbx, [g]`
`mov rax, 1`
`add rax, [rbx]`
`mov [rbx], rax`
`xor rbx, rbx`

`rax=?, rbx=?, [g]=3`

`rax=?, rbx=@, [g]=3`

`rax=?, rbx=@, [g]=3`

`rax=34?, rbx=g, [g]=3`

`rax=3, rbx=?, [g]=3`

`rax=3, rbx=0, [g]=3`

```
lea rbx, [g]
```

```
mov rax, 1
```

```
add rax, [rbx]
```

```
➡ mov [rbx], rax
```

```
xor rbx, rbx
```

rax=?, rbx=?, [g]=?

rax=?, rbx=g, [g]=?

rax=1, rbx=g, [g]=?

rax=3, rbx=g, [g]=?

rax=3, rbx=@, [g]=3

rax=3, rbx=0, [g]=3

lea rbx, [g]

mov rax, 1

add rax, [rbx]

mov [rbx], rax

➡ xor rbx, rbx

rax=?, rbx=?, [g]=2

rax=?, rbx=g, [g]=2

rax=1, rbx=g, [g]=2?

rax=3, rbx=g, [g]=?

rax=3, rbx=g, [g]=3

rax=3, rbx=0, [g]=3

➔ `lea rbx, [g]`
`mov rax, 1`
`add rax, [rbx]`
`mov [rbx], rax`
`xor rbx, rbx`

`rax=?, rbx=?, [g]=2`

`rax=?, rbx=g, [g]=2`

`rax=1, rbx=g, [g]=2`

`rax=3, rbx=g, [g]=2`

`rax=3, rbx=g, [g]=3`

`rax=3, rbx=0, [g]=3`

lea rbx, [g]

mov rax, 1

add rax, [rbx]

➡ mov [rbx], rax

xor rbx, rbx

rax=?, rbx=?, [g]=2

rax=?, rbx=g, [g]=2

rax=1, rbx=g, [g]=2

rax=3, rbx=g, [g]=2

rax=3, rbx=g, [g]=3

rax=3, rbx=0, [g]=3

Key Techniques

- Forward Execution
 - Recovers states before irreversible instructions
- Error Correction
 - Handles errors introduced by “missing” memory writes




REPT Data Recovery

- Single-threaded execution reconstruction
- Multi-threaded execution reconstruction


Core Dump

```
10101011010101010101011001101010010011110110110
11010101101010101101010101010110011010100100
1111011011010101011010101101001010101010110
011010100100111101101101010101101010110101010
10101011001101010010011110110110101011011010
1011010010101010101011001101010010011110110110
10101101101010110101010101010110011010100100111
10110110110101101101010101010101010101100110
10100100111101101101010110110101101010110101
01010101011001101010010011110110110101011011
01010110101010101010110011010100100111101101
01011010101010101010011010100100111101101101
101011011010110101010101010110011010100100111
10110110101011010101010101011010101010101010
11001101010010011110110110101011011010101101
01010101011001101010010011110110110101010101
0101010110011010100100111101101101010101101
0101101010101010101011001101010010011110110110
10101101101011010101010101010101010101101010
01001111011011010101011011010110101010101010
110011010100100111101101101010110110101010101
0110101010101010110011010100100111101101101010
```

Instruction Sequence #1



Instruction Sequence #2






Execution History

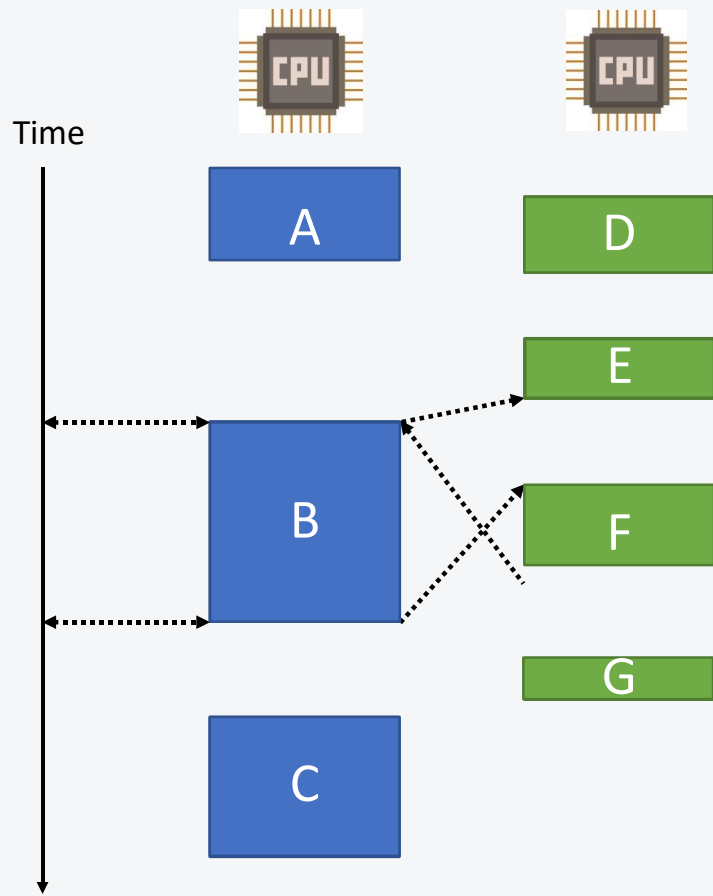
= ?





How to determine the thread interleavings?

Time

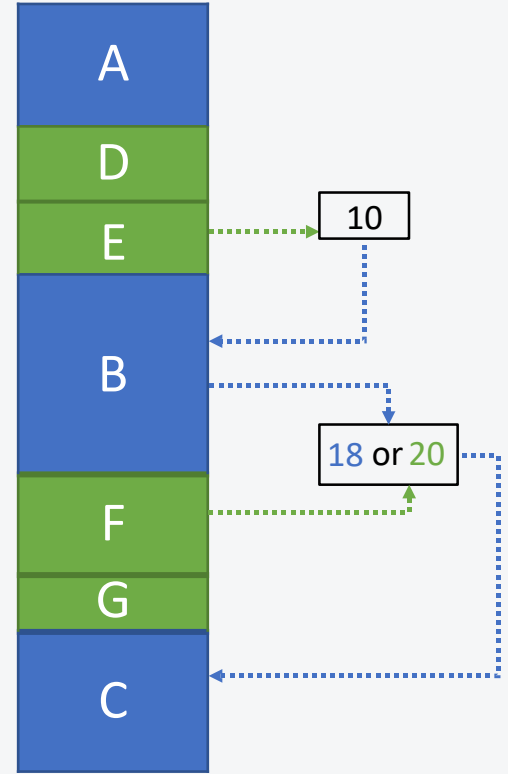







Time



Time



Time




Key Techniques

- Hardware Timestamps
 - Constructs a partial order
- Concurrent memory write detection
 - Constrains their usage to avoid propagating a wrong value



With REPT



Evaluation




16 bugs



1-5% overhead



14 bugs



92% accuracy

Limitations

- The control flow trace may not be long enough to capture the defect e.g., the free call is not in the trace for a use-after-free bug.
- Data values that are necessary for debugging the failure are not recovered e.g., the heap address passed to the free call is not recovered for a use-after-free bug.

Future Work

- REPT currently does not capture any data during a program's execution. To fundamentally solve these two limitations, we will need to log more data than just the memory dump.
- While REPT's core analysis is on machine instructions and thus independent of the privilege mode, we need to properly handle kernel-specific artifacts such as interrupts to support reverse debugging of kernel-mode executions.
- Can study how program tracing can be combined with event logging to help developers debug bugs in distributed systems.

Conclusion

- Debugging production failures is important but hard
- REPT is a practical reverse debugging solution for production failures
 - Online hardware tracing to log the control flow with timestamps
 - Offline binary analysis to recover the data flow with high accuracy
- REPT has been deployed on Microsoft Windows

REFERENCES



Cui, Weidong, et al. "{REPT}: Reverse debugging of failures in deployed software." *13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18)*. 2018.



Cui, Weidong, et al. "Retracer: Triaging crashes by reverse execution from partial memory dumps." *2016 IEEE/ACM 38th International Conference on Software Engineering (ICSE)*. IEEE, 2016.



THANKYOU